

CSE527 Homework 4

Due date: 23:59 on Nov. 5, 2019 (Tuesday)

In this semester, we will use Google Colab for the assignments, which allows us to utilize resources that some of us might not have in their local machines such as GPUs. You will need to use your Stony Brook (*.stonybrook.edu) account for coding and Google Drive to save your results.

Google Colab Tutorial

Go to <https://colab.research.google.com/notebooks/>, you will see a tutorial named "Welcome to Colaboratory" file, where you can learn the basics of using google colab.

Settings used for assignments: **Edit -> Notebook Settings -> Runtime Type (Python 3)**.

Description

This project is an introduction to deep learning tools for computer vision. You will design and train deep convolutional networks for scene recognition using [PyTorch](#). You can visualize the structure of the network with [mNeuron] (http://vision03.csail.mit.edu/cnn_art/index.html)

Remember Homework 3: Scene recognition with bag of words. You worked hard to design a bag of features representations that achieved 60% to 70% accuracy (most likely) on 16-way scene classification. We're going to attack the same task with deep learning and get higher accuracy. Training from scratch won't work quite as well as homework 3 due to the insufficient amount of data, fine-tuning an existing network will work much better than homework 3.

In Problem 1 of the project you will train a deep convolutional network from scratch to recognize scenes. The starter codes gives you methods to load data and display them. You will need to define a simple network architecture and add jittering, normalization, and regularization to increase recognition accuracy to 50, 60, or perhaps 70%. Unfortunately, we only have 2,400 training examples so it doesn't seem possible to train a network from scratch which outperforms hand-crafted features

For Problem 2 you will instead fine-tune a pre-trained deep network to achieve about 85% accuracy on the task. We will use the pretrained AlexNet network which was not trained to recognize scenes at all.

These two approaches represent the most common approaches to recognition problems in computer vision today -- train a deep network from scratch if you have enough data (it's not always obvious whether or not you do), and if you cannot then instead fine-tune a pre-trained network.

There are 2 problems in this homework with a total of 110 points including 10 bonus points. Be sure to read **Submission Guidelines** below. They are important. For the problems requiring text descriptions, you might want to add a markdown block for that.

Dataset

Save the [dataset\(click me\)](#) into your working folder in your Google Drive for this homework.
Under your root folder, there should be a folder named "data" (i.e. XXX/Surname_Givenname_SBUID/data) containing the images. **Do not upload**

the data subfolder before submitting on blackboard due to size limit. There should be only one .ipynb file under your root folder Surname_Givenname_SBUID.

Some Tutorials (PyTorch)

- You will be using PyTorch for deep learning toolbox (follow the [link](#) for installation).
- For PyTorch beginners, please read this [tutorial](#) before doing your homework.
- Feel free to study more tutorials at <http://pytorch.org/tutorials/>.
- Find cool visualization here at <http://playground.tensorflow.org>.

Starter Code

```
1  # import packages here
2  import cv2
3  import numpy as np
4  import matplotlib.pyplot as plt
5  import glob
6  import random
7  import time
8
9  # for creating validation set
10 from sklearn.model_selection import train_test_split
11
12 # for evaluating the model
13 from sklearn.metrics import accuracy_score
14 from tqdm import tqdm
15
16 import torch
17 import torchvision
18 from torchvision import models
19 import torchvision.transforms as transforms
20
21 from torch.autograd import Variable
22 import torch.nn as nn
23 import torch.nn.functional as F
24
25 # PyTorch libraries and modules
26 from torch.nn import Linear, ReLU, CrossEntropyLoss, Sequential, Conv2d, MaxPool2d, Module, Softmax, BatchNorm2d, Dropout
27 from torch.optim import Adam, SGD
28
29 import pickle
30 import time
```



```
1  # Mount your google drive where you've saved your assignment folder
2  from google.colab import drive
3  drive.mount('/content/gdrive', force_remount=True)
4  %pwd
```



Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Awg%3Aoauth%3A2.0%3Aoob&s

Enter your authorization code:
.....
Mounted at /content/gdrive
'/content'

```
1 # Set your working directory (in your google drive)
2 # Note that 'gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/hw4' is just an example,
3 #   change it to your specific homework directory.
4 cd '/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw4'
```

```
📁 /content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw4
```

```
1 %pwd
2 %ls -lrt
```

```
📁 total 2007647
-rw----- 1 root root    39371054 Nov  4 01:33 trainloader_small.pkl
-rw----- 1 root root     6561642 Nov  4 01:33 testloader_small.pkl
-rw----- 1 root root    78742766 Nov  4 06:28 trainloader_small_augmented.pkl
-rw----- 1 root root    78742766 Nov  4 06:37 trainloader_small_augmented_normalized.pkl
-rw----- 1 root root   118113902 Nov  4 08:59 trainloader_small_augmented_normalized_rotated.pkl
-rw----- 1 root root   1445118158 Nov  4 09:27 trainloader_large.pkl
-rw----- 1 root root    240852826 Nov  4 09:30 testloader_large.pkl
-rw----- 1 root root    39336214 Nov  5 02:08 alexnet_output.pkl
-rw----- 1 root root     6556058 Nov  5 02:08 alexnet_output_test.pkl
-rw----- 1 root root     2346954 Nov  5 18:16 Paratkar_Shreyash_112673930_hw4.pdf
-rw----- 1 root root         85678 Nov  6 00:14 Paratkar_Shreyash_112673930_hw4.ipynb
```

```
1 # =====
2 #   Load Training Data and Testing Data
3 # =====
4 class_names = [name[111:] for name in glob.glob('/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw3/data/train/*')]
5 class_names = dict(zip(range(len(class_names)), class_names))
6 print("class_names: %s " % class_names)
7 n_train_samples = 150
8 n_test_samples = 50
9
10
11 def img_norm(img):
12     #
13     # Write your code here
14     # normalize img pixels to [-1, 1]
15     #
16     img = img.astype(float)
17     img = cv2.normalize(img, None, alpha=-1, beta=1, norm_type=cv2.NORM_MINMAX, dtype=cv2.CV_32F)
18     return img
19
20 def load_dataset(path, img_size, num_per_class=-1, batch_num=1, shuffle=False, augment=False, is_color=False,
21                 rotate_90=False, zero_centered=False):
22
```

```

23     data = []
24     labels = []
25
26     if is_color:
27         channel_num = 3
28     else:
29         channel_num = 1
30
31     # read images and resizing
32     for id, class_name in class_names.items():
33         print("Loading images from class: %s" % id)
34         img_path_class = glob.glob(path + class_name + '/*.jpg')
35         if num_per_class > 0:
36             img_path_class = img_path_class[:num_per_class]
37         labels.extend([id]*len(img_path_class))
38         if augment == True:
39             labels.extend([id]*len(img_path_class))
40         if rotate_90 == True:
41             labels.extend([id]*len(img_path_class))
42         for filename in img_path_class:
43             if is_color:
44                 img = cv2.imread(filename)
45             else:
46                 img = cv2.imread(filename, 0)
47
48             # resize the image
49             img = cv2.resize(img, img_size, cv2.INTER_LINEAR)
50
51             if is_color:
52                 img = np.transpose(img, [2, 0, 1])
53
54             # norm pixel values to [-1, 1]
55             data.append(img_norm(img))
56             if augment == True:
57                 data.append(np.fliplr(img_norm(img)))
58             if rotate_90 == True:
59                 data.append(np.rot90(img_norm(img)))
60
61     #
62     # Write your Data Augmentation code here
63     # mirroring
64     #
65
66     #
67     # Write your Data Normalization code here
68     # norm data to zero-centered
69     #
70
71     if zero_centered == True:
72         data = data - np.mean(data)
73
74     # randomly permute (this step is important for training)
75     if shuffle:
76         bundle = list(zip(data, labels))

```

```

76         bundle = list(zip(data, labels))
77         random.shuffle(bundle)
78         random.shuffle(bundle)
79         data, labels = zip(*bundle)
80
81     # divide data into minibatches of TorchTensors
82     if batch_num > 1:
83         batch_data = []
84         batch_labels = []
85
86         print(len(data))
87         print(batch_num)
88
89         for i in range(int(len(data) / batch_num)):
90             minibatch_d = data[i*batch_num: (i+1)*batch_num]
91             minibatch_d = np.reshape(minibatch_d, (batch_num, channel_num, img_size[0], img_size[1]))
92             batch_data.append(torch.from_numpy(minibatch_d))
93
94             minibatch_l = labels[i*batch_num: (i+1)*batch_num]
95             batch_labels.append(torch.LongTensor(minibatch_l))
96         data, labels = batch_data, batch_labels
97
98     return zip(batch_data, batch_labels)

```


class_names: {0: 'Forest', 1: 'Industrial', 2: 'Flower', 3: 'Coast', 4: 'InsideCity', 5: 'Office', 6: 'Bedroom', 7: 'Highway', 8: 'Street', 9: 'TallBuilding', 10: 'LivingRoom', 11: 'Suburb', 12: 'OpenCoun

```

1  # load data into size (64, 64)
2  img_size = (64, 64)
3  batch_num = 50 # training sample number per batch
4
5  # load training dataset
6  trainloader_small = list(load_dataset('/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw3/data/train/', img_size, batch_num=batch_num, shuffle=True, augmen
7  train_num = len(trainloader_small)
8  print("Finish loading %d minibatches(=%d) of training samples." % (train_num, batch_num))
9
10 # load testing dataset
11 testloader_small = list(load_dataset('/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw3/data/test/', img_size, num_per_class=50, batch_num=batch_num))
12 test_num = len(testloader_small)
13 print("Finish loading %d minibatches(=%d) of testing samples." % (test_num, batch_num))

```



```
Loading images from class: 0
Loading images from class: 1
Loading images from class: 2
Loading images from class: 3
Loading images from class: 4
Loading images from class: 5
Loading images from class: 6
Loading images from class: 7
Loading images from class: 8
Loading images from class: 9
Loading images from class: 10
Loading images from class: 11
Loading images from class: 12
Loading images from class: 13
Loading images from class: 14
Loading images from class: 15
2400
50
Finish loading 48 minibatches(=50) of training samples.
Loading images from class: 0
Loading images from class: 1
Loading images from class: 2
Loading images from class: 3
Loading images from class: 4
Loading images from class: 5
Loading images from class: 6
Loading images from class: 7
Loading images from class: 8
Loading images from class: 9
Loading images from class: 10
Loading images from class: 11
Loading images from class: 12
Loading images from class: 13
Loading images from class: 14
Loading images from class: 15
400
50
Finish loading 8 minibatches(=50) of testing samples.
```

```
1  file = open('trainloader_small.pkl','wb')
2  pickle.dump(trainloader_small, file)
3  file.close()
4  file = open('testloader_small.pkl','wb')
5  pickle.dump(testloader_small, file)
6  file.close()
```

```
1  # Reading from pickle files
2  file = open('trainloader_small.pkl', 'rb')
3  trainloader_small = pickle.load(file)
4  file.close()
5  file = open('testloader_small.pkl', 'rb')
6  testloader_small = pickle.load(file)
7  file.close()
```


```
1  type(trainloader_small[47][0][49])
```

```
2 # trainloader_small[47][0][49].shape -- image torch {47th batch , 0 for selecting image, 49th image in batch of 50}
3 # trainloader_small[47][1][49] -- label {47th batch , 1 for selecting label, 49th image in batch of 50}
4 # trainloader_small[batch_index][img/label][50 images][0]
5 # for i in range(48):
6 #     print(trainloader_small[i][0].shape, trainloader_small[i][1].shape)
```

 torch.Tensor

```
1 # show some images
2 def imshow(img):
3     img = img / 2 + 0.5      # unnormalize
4     npimg = img.numpy()
5     if len(npimg.shape) > 2:
6         npimg = np.transpose(img, [1, 2, 0])
7     plt.figure
8     plt.imshow(npimg, 'gray')
9     plt.show()
10 img, label = trainloader_small[0][0][11][0], trainloader_small[0][1][11]
11 label = int(np.array(label))
12 print(class_names[label])
13 imshow(img)
```

 Store



▼ Problem 1: Training a Network From Scratch

{Part 1: 35 points} Gone are the days of hand designed features. Now we have end-to-end learning in which a highly non-linear representation is learned for our data to maximize our objective (in this case, 16-way classification accuracy). Instead of 70% accuracy we can now recognize scenes with... 25% accuracy. OK, that didn't work at all. Try to boost the accuracy by doing the following:

Data Augmentation: We don't have enough training data, let's augment the training data. If you left-right flip (mirror) an image of a scene, it never changes categories. A kitchen doesn't become a forest when mirrored. This isn't true in all domains — a "d" becomes a "b" when mirrored, so you can't "jitter" digit recognition training data in the same way. But we can synthetically increase our amount of training data by left-right mirroring training images during the learning process.

After you implement mirroring, you should notice that your training error doesn't drop as quickly. That's actually a good thing, because it means the network isn't overfitting to the 2,400 original training images as much (because it sees 4,800 training images now, although they're not as

good as 4,800 truly independent samples). Because the training and test errors fall more slowly, you may need more training epochs or you may try modifying the learning rate. You should see a roughly 10% increase in accuracy by adding mirroring. You are **required** to implement mirroring as data augmentation for this part.

You can try more elaborate forms of jittering -- zooming in a random amount, rotating a random amount, taking a random crop, etc. These are not required, you might want to try these in the bonus part.

Data Normalization: The images aren't zero-centered. One simple trick which can help a lot is to subtract the mean from every image. It would arguably be more proper to only compute the mean from the training images (since the test/validation images should be strictly held out) but it won't make much of a difference. After doing this you should see another 15% or so increase in accuracy. This part is **required**.

Network Regularization: Add dropout layer. If you train your network (especially for more than the default 30 epochs) you'll see that the training error can decrease to zero while the val top1 error hovers at 40% to 50%. The network has learned weights which can perfectly recognize the training data, but those weights don't generalize to held out test data. The best regularization would be more training data but we don't have that. Instead we will use dropout regularization.

What does dropout regularization do? It randomly turns off network connections at training time to fight overfitting. This prevents a unit in one layer from relying too strongly on a single unit in the previous layer. Dropout regularization can be interpreted as simultaneously training many "thinned" versions of your network. At test, all connections are restored which is analogous to taking an average prediction over all of the "thinned" networks. You can see a more complete discussion of dropout regularization in this [paper](#).

The dropout layer has only one free parameter — the dropout rate — the proportion of connections that are randomly deleted. The default of 0.5 should be fine. Insert a dropout layer between your convolutional layers. In particular, insert it directly before your last convolutional layer. Your test accuracy should increase by another 10%. Your train accuracy should decrease much more slowly. That's to be expected — you're making life much harder for the training algorithm by cutting out connections randomly.

If you increase the number of training epochs (and maybe decrease the learning rate) you should be able to achieve around 50% test accuracy. In this part, you are **required** to add dropout layer to your network.

Please give detailed descriptions of your network layout in the following format:

Data augmentation: [descriptions]

Data normalization: [descriptions]

Layer 1: [layer_type]: [Parameters]

Layer 2: [layer_type]: [Parameters]

...

Then report the final accuracy on test set and time consumed for training and testing separately.

{Part 2: 15 points} Try **three techniques** taught in the class to increase the accuracy of your model. Such as increasing training data by randomly rotating training images, adding batch normalization, different activation functions (e.g., sigmoid) and model architecture modification. Note that too many layers can do you no good due to insufficient training data. Clearly describe your method and accuracy increase/decrease for each of the three techniques.

```
1  # =====
2  #      Define Network Architecture
3  # =====
4  class Net(Module):
5      def __init__(self):
```



```
5     def __init__(self):
6         super(Net, self).__init__()
7
8         self.cnn_layers = Sequential(
9             Conv2d(1, 32, kernel_size=5, stride=1, padding=2),
10             ReLU(),
11             MaxPool2d(kernel_size=2, stride=2),
12             Conv2d(32, 64, kernel_size=5, stride=1, padding=2),
13             ReLU(),
14             MaxPool2d(kernel_size=2, stride=2),
15         )
16         self.linear_layers = Sequential(
17             Linear(64 * 16 * 16, 1000),
18             Linear(1000, 16)
19         )
20     def forward(self, x):
21         x = self.cnn_layers(x)
22         x = x.view(x.size(0), -1)
23         x = self.linear_layers(x)
24         return x
25
26 class Net_Dropout(Module):
27     def __init__(self):
28         super(Net_Dropout, self).__init__()
29         self.layer1 = nn.Sequential(
30             nn.Conv2d(1, 32, kernel_size=5, stride=1, padding=2),
31             nn.ReLU(),
32             nn.MaxPool2d(kernel_size=2, stride=2))
33         self.layer2 = nn.Sequential(
34             nn.Conv2d(32, 64, kernel_size=5, stride=1, padding=2),
35             nn.ReLU(),
36             nn.MaxPool2d(kernel_size=2, stride=2))
37         self.drop_out = nn.Dropout(p=0.5)
38         self.fc1 = nn.Linear(16 * 16 * 64, 1000)
39         self.fc2 = nn.Linear(1000, 16)
40
41     def forward(self, x):
42         out = self.layer1(x)
43         out = self.drop_out(out)
44         out = self.layer2(out)
45         out = out.reshape(out.size(0), -1)
46         out = self.fc1(out)
47         out = self.fc2(out)
48         return out
49
50 class Net_Enhance1(Module):
51     def __init__(self):
52         super(Net_Enhance1, self).__init__()
53         self.layer1 = nn.Sequential(
54             nn.Conv2d(1, 32, kernel_size=5, stride=1, padding=2),
55             BatchNorm2d(32),
56             nn.ReLU(),
57             nn.MaxPool2d(kernel_size=2, stride=2))
58         self.layer2 = nn.Sequential(
```

```

59         nn.Conv2d(32, 64, kernel_size=5, stride=1, padding=2),
60         BatchNorm2d(64),
61         nn.ReLU(),
62         nn.MaxPool2d(kernel_size=2, stride=2))
63     self.drop_out = nn.Dropout(p=0.5)
64     self.fc1 = nn.Linear(16 * 16 * 64, 1000)
65     self.fc2 = nn.Linear(1000, 16)
66
67     def forward(self, x):
68         out = self.layer1(x)
69         out = self.drop_out(out)
70         out = self.layer2(out)
71         out = out.reshape(out.size(0), -1)
72         out = self.fc1(out)
73         out = self.fc2(out)
74         return out
75
76 class Net_Enhance2(Module):
77     def __init__(self):
78         super(Net_Enhance2, self).__init__()
79         self.layer1 = nn.Sequential(
80             nn.Conv2d(1, 32, kernel_size=5, stride=1, padding=2),
81             nn.Sigmoid(),
82             nn.MaxPool2d(kernel_size=2, stride=2))
83         self.layer2 = nn.Sequential(
84             nn.Conv2d(32, 64, kernel_size=5, stride=1, padding=2),
85             nn.Sigmoid(),
86             nn.MaxPool2d(kernel_size=2, stride=2))
87         self.drop_out = nn.Dropout(p=0.5)
88         self.fc1 = nn.Linear(16 * 16 * 64, 1000)
89         self.fc2 = nn.Linear(1000, 16)
90
91     def forward(self, x):
92         out = self.layer1(x)
93         out = self.drop_out(out)
94         out = self.layer2(out)
95         out = out.reshape(out.size(0), -1)
96         out = self.fc1(out)
97         out = self.fc2(out)
98         return out
99
100 def train(epoch, training_set, device, model, optimizer, criterion):
101     for i in range(len(training_set)):
102         x_train, y_train = Variable(training_set[i][0]).cuda(device), Variable(training_set[i][1]).cuda(device)
103         output_train = model(x_train)
104         loss_train = criterion(output_train, y_train)
105         optimizer.zero_grad()
106         loss_train.backward()
107         optimizer.step()
108
109 def test(training_set, device, model):
110     with torch.no_grad():
111         total = 0
112         correct = 0

```

```
112         correct += 1
113     for i in range(8):
114         outputs = model(Variable(training_set[i][0]).cuda(device))
115         _, predicted = torch.max(outputs.data, 1)
116         total += training_set[i][1].size(0)
117         correct += (predicted == training_set[i][1].cuda(device)).sum().item()
118     print('Test Accuracy of the model on test images: {} %'.format((correct / total) * 100))
```

```
1  device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
2  model = Net()
3  model.cuda(device)
4  optimizer = Adam(model.parameters(), lr=0.001)
5  criterion = CrossEntropyLoss()
6  n_epochs = 5
7  start_time = time.time()
8  for epoch in range(n_epochs):
9      train(epoch, trainloader_small, device, model, optimizer, criterion)
10 print("Time taken for training: ", round(time.time() - start_time), "seconds")
11 start_time = time.time()
12 test(testloader_small, device, model)
13 print("Time taken for testing: ", round(time.time() - start_time), "seconds")
```

```
📄 Time taken for training:  2 seconds
Test Accuracy of the model on test images: 24.5 %
Time taken for testing:  0 seconds
```

Test Accuracy acheived with the basic architecture and training size of **2400** images = **24.5%**

```
1  # =====
2  #          Optimize/Train Network
3  # =====
4  trainloader_small_augmented = list(load_dataset('/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw3/data/train/', img_size, batch_num=batch_num, shuffle=True))
5  train_num_augmented = len(trainloader_small_augmented)
6  print("Finish loading %d minibatches(=%d) of training samples." % (train_num_augmented, batch_num))
7  file = open('trainloader_small_augmented.pkl','wb')
8  pickle.dump(trainloader_small_augmented, file)
9  file.close()
```

```
📄
```

```
Loading images from class: 0
Loading images from class: 1
Loading images from class: 2
Loading images from class: 3
Loading images from class: 4
Loading images from class: 5
Loading images from class: 6
Loading images from class: 7
Loading images from class: 8
Loading images from class: 9
Loading images from class: 10
Loading images from class: 11
Loading images from class: 12
Loading images from class: 13
Loading images from class: 14
Loading images from class: 15
4800
50
Finish loading 96 minibatches(=50) of training samples.
```

Data Augmented to contain **2400** more images which are the mirrored/laterally inverted versions of the original **2400** images, thereby giving us a final images set of **4800**.

```
1  file = open('trainloader_small_augmented.pkl', 'rb')
2  trainloader_small_augmented = pickle.load(file)
3  file.close()

1  device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
2  model = Net()
3  model.cuda(device)
4  optimizer = Adam(model.parameters(), lr=0.001)
5  criterion = CrossEntropyLoss()
6  start_time = time.time()
7  for epoch in range(n_epochs):
8      train(epoch, trainloader_small_augmented, device, model, optimizer, criterion)
9  print("Time taken for training: ", round(time.time() - start_time), "seconds")
10 start_time = time.time()
11 test(testloader_small, device, model)
12 print("Time taken for testing: ", round(time.time() - start_time), "seconds")
```

```
📄 Time taken for training:  5 seconds
Test Accuracy of the model on test images: 34.25 %
Time taken for testing:  0 seconds
```

The accuracy obtained after **Data Augmentation** = **34.25%**

```
1  trainloader_small_augmented_normalized = list(load_dataset('/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw3/data/train/', img_size, batch_num=batch_num,
2  train_num_augmented_normalized = len(trainloader_small_augmented_normalized)
3  print("Finish loading %d minibatches(=%d) of training samples." % (train_num_augmented_normalized, batch_num))
4  file = open('trainloader_small_augmented_normalized.pkl','wb')
5  pickle.dump(trainloader_small_augmented_normalized, file)
6  file.close()
```

```
0 file.close()
```

```
↳ Loading images from class: 0
Loading images from class: 1
Loading images from class: 2
Loading images from class: 3
Loading images from class: 4
Loading images from class: 5
Loading images from class: 6
Loading images from class: 7
Loading images from class: 8
Loading images from class: 9
Loading images from class: 10
Loading images from class: 11
Loading images from class: 12
Loading images from class: 13
Loading images from class: 14
Loading images from class: 15
4800
50
Finish loading 96 minibatches(=50) of training samples.
```

Images are **zero-centered** by subtracting each pixel of each image by the mean of all the pixels in the dataset.

```
1 file = open('trainloader_small_augmented_normalized.pkl', 'rb')
2 trainloader_small_augmented_normalized = pickle.load(file)
3 file.close()

1 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
2 model = Net()
3 model.cuda(device)
4 optimizer = Adam(model.parameters(), lr=0.001)
5 criterion = CrossEntropyLoss()
6 start_time = time.time()
7 for epoch in range(n_epochs):
8     train(epoch, trainloader_small_augmented_normalized, device, model, optimizer, criterion)
9 print("Time taken for training: ", round(time.time() - start_time), "seconds")
10 start_time = time.time()
11 test(testloader_small, device, model)
12 print("Time taken for testing: ", round(time.time() - start_time), "seconds")
```

```
↳ Time taken for training:  5 seconds
Test Accuracy of the model on test images: 41.25 %
Time taken for testing:  0 seconds
```

Model accuracy after **Data Normalization = 41.25%**

```
1 model = Net_Dropout()
2 model.cuda(device)
3 optimizer = Adam(model.parameters(), lr=0.001)
4 criterion = CrossEntropyLoss()
5 start_time = time.time()
6 for epoch in range(20):
7     train(epoch, trainloader_small augmented normalized, device, model, optimizer, criterion)
```

```
8 print("Time taken for training: ", round(time.time() - start_time), "seconds")
9 start_time = time.time()
10 test(testloader_small, device, model)
11 print("Time taken for testing: ", round(time.time() - start_time), "seconds")
```

```
➤ Time taken for training: 19 seconds
  Test Accuracy of the model on test images: 47.0 %
  Time taken for testing: 0 seconds
```

A dropout layer is added before the 2nd convolutional layer, with a default droupout rate of **0.5**. Accuracy obtained after adding dropout = **47%**

Data augmentation: To increase the accuracy of the convolutional neural network, we augment the training dataset with the laterally inverted variants of the same original images, since the labels for these images would remain the same even if the orientation changes.

This augmentation gives us an accuracy of **34.25%**, which is a development on the original **24.5%**.

The time taken for training is 5 seconds while the testing gets completed in ~0 seconds.

Data normalization: To zero-center te images, the mean of the whole dataset's images, is subtracted from all the images, as the normalization step. This gives our model an accuracy of **41.25%** and takes about the same time for training and testing as the previous method, i.e. 5 seconds and ~0 seconds respectively.

The convolutional neural network defined for the basic test, data augmentation and data normalization involved the following layers.

Layer 1: Sequential layer: Consists of convolution + ReLU + pooling sequences.

i.) Conv2d nn.Module creates a set of convolutional filters : [1 input channel(grayscale image), 32 channels for first convolutional filter layer, kernel_size = 5, stride = 1, padding = 2]

ii.) A simple ReLU activation.

iii.) Max pooling operation, pooling size = 2. To downsample by size of 2, we set the stride = 2 and padding = 0 -> 32 channels of 32 x 32 “images”.

Layer 2: Sequential layer: Consists of convolution + ReLU + pooling sequences.

i.) Conv2d nn.Module creates a set of convolutional filters : [32 input channel, 64 channels for second convolutional filter layer, kernel_size = 5, stride = 1, padding = 2]

ii.) A simple ReLU activation.

iii.) Max pooling operation, pooling size = 2. To downsample by size of 2, we set the stride = 2 and padding = 0 -> 32 channels of 16 x 16 “images”.

Layer 3: Fully Connected Layer: The first layer will be of size 16 x 16 x 64 nodes.

Layer 4: Fully Connected Layer: first layer will connect to the second layer of 1000 nodes.

Layer 1.5: Dropout Layer: A dropout layer is inserted as mentioned in the question before the second convolutional layer.

The basic CNN gives an accuracy of **24.5%** with training time of **5** seconds.

The CNN implemented **with dropout** gives a testing accuracy of **47%** with training time of **19** seconds with **20 epochs**.

All the test accuracy values measured fluctuate above and below the above shown values.

```
1  # =====
2  #           Evaluating Network
3  # =====
4  model = Net_Enhance1()
5  model.cuda(device)
6  optimizer = Adam(model.parameters(), lr=0.001)
7  criterion = CrossEntropyLoss()
8  start_time = time.time()
9  for epoch in range(n_epochs):
10     train(epoch, trainloader_small_augmented_normalized, device, model, optimizer, criterion)
11 print("Time taken for training: ", round(time.time() - start_time), "seconds")
12 start_time = time.time()
13 test(testloader_small, device, model)
14 print("Time taken for testing: ", round(time.time() - start_time), "seconds")
```

Time taken for training: 5 seconds
Test Accuracy of the model on test images: 32.5 %
Time taken for testing: 0 seconds

```
1  model = Net_Enhance2()
2  model.cuda(device)
3  optimizer = Adam(model.parameters(), lr=0.001)
4  criterion = CrossEntropyLoss()
5  start_time = time.time()
6  for epoch in range(n_epochs):
7     train(epoch, trainloader_small_augmented_normalized, device, model, optimizer, criterion)
8  print("Time taken for training: ", round(time.time() - start_time), "seconds")
9  start_time = time.time()
10 test(testloader_small, device, model)
11 print("Time taken for testing: ", round(time.time() - start_time), "seconds")
```

Time taken for training: 5 seconds
Test Accuracy of the model on test images: 25.25 %
Time taken for testing: 0 seconds

```
1  img_size = (64, 64)
2  batch_num = 50
3  trainloader_small_augmented_normalized_rotated = list(load_dataset('/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw3/data/train/', img_size, batch_num=batch_num))
4  train_num_small_augmented_normalized_rotated = len(trainloader_small_augmented_normalized_rotated)
5  print("Finish loading %d minibatches(=%d) of training samples." % (train_num_small_augmented_normalized_rotated, batch_num))
6  file = open('trainloader_small_augmented_normalized_rotated.pkl', 'wb')
7  pickle.dump(trainloader_small_augmented_normalized_rotated, file)
8  file.close()
```

```
Loading images from class: 0
Loading images from class: 1
Loading images from class: 2
Loading images from class: 3
Loading images from class: 4
Loading images from class: 5
Loading images from class: 6
Loading images from class: 7
Loading images from class: 8
Loading images from class: 9
Loading images from class: 10
Loading images from class: 11
Loading images from class: 12
Loading images from class: 13
Loading images from class: 14
Loading images from class: 15
7200
50
Finish loading 144 minibatches(=50) of training samples.
```

```
1  file = open('trainloader_small_augmented_normalized_rotated.pkl', 'rb')
2  trainloader_small_augmented_normalized_rotated = pickle.load(file)
3  file.close()

1  model = Net_Dropout()
2  model.cuda(device)
3  optimizer = Adam(model.parameters(), lr=0.001)
4  criterion = CrossEntropyLoss()
5  start_time = time.time()
6  for epoch in range(20):
7      train(epoch, trainloader_small_augmented_normalized_rotated, device, model, optimizer, criterion)
8  print("Time taken for training: ", round(time.time() - start_time), "seconds")
9  start_time = time.time()
10 test(testloader_small, device, model)
11 print("Time taken for testing: ", round(time.time() - start_time), "seconds")
```

```
📄➡ Time taken for training:  28 seconds
Test Accuracy of the model on test images: 48.75 %
Time taken for testing:  0 seconds
```

Three different techniques were tried to increase the accuracy of the baseline model which is trained for 3 epochs.

The first two techniques were passed augmented and normalized input sets of 4800 images.

1.) **Adding Batch Normalization:** A batch normalization element was added to the sequential layers of the basic CNN constructed. It was added twice after each of the convolution elements of the repective layers and was passed the number of output channels of the convolution elements as input (32 and 64 respectively).

This resulted in a test accuracy of **32.5%**, which is a decrease as compared to the **24.5%** obtained by the original network. Time taken for training: **5** seconds.

2.) **Using sigmoid activation function instead of ReLU:** Using the sigmoid activation function instead of ReLU gave a **0.75%** increase in test accuracy over the one obtained by using ReLU.

Test accuracy with sigmoid activation function = **25.25%**. Time taken for training: **5** seconds.

3.) For the third attempt at increasing accuracy, the previously increased training dataset of 4800 images was further augmented to contain 2400 more images which are the original images rotated by 90 degrees each.

This resulted in an increase of **24.25%** in test accuracy over the original model.

Test accuracy with sigmoid activation function = **48.75%**. Because the training data size is tripled, the training time required with 20 epochs is

Resource referred:

1.) <https://www.analyticsvidhya.com/blog/2019/10/building-image-classification-models-cnn-pytorch/>

2.) <https://adventuresinmachinelearning.com/convolutional-neural-networks-tutorial-in-pytorch/>

▼ Problem 2: Fine Tuning a Pre-Trained Deep Network

{Part 1: 30 points} Our convolutional network to this point isn't "deep". Fortunately, the representations learned by deep convolutional networks is that they generalize surprisingly well to other recognition tasks.

But how do we use an existing deep network for a new recognition task? Take for instance, [AlexNet](#) network has 1000 units in the final layer corresponding to 1000 ImageNet categories.

Strategy A: One could use those 1000 activations as a feature in place of a hand crafted feature such as a bag-of-features representation. You would train a classifier (typically a linear SVM) in that 1000 dimensional feature space. However, those activations are clearly very object specific and may not generalize well to new recognition tasks. It is generally better to use the activations in slightly earlier layers of the network, e.g. the 4096 activations in the last 2nd fully-connected layer. You can often get away with sub-sampling those 4096 activations considerably, e.g. taking only the first 200 activations.

Strategy B: *Fine-tune* an existing network. In this scenario you take an existing network, replace the final layer (or more) with random weights, and train the entire network again with images and ground truth labels for your recognition task. You are effectively treating the pre-trained deep network as a better initialization than the random weights used when training from scratch. When you don't have enough training data to train a complex network from scratch (e.g. with the 16 classes) this is an attractive option. Fine-tuning can work far better than Strategy A of taking the activations directly from an pre-trained CNN. For example, in [this paper](#) from CVPR 2015, there wasn't enough data to train a deep network from scratch, but fine tuning led to 4 times higher accuracy than using off-the-shelf networks directly.

You are required to implement **Strategy B** to fine-tune a pre-trained **AlexNet** for this scene classification task. You should be able to achieve performance of 85% approximately. It takes roughly 35~40 minutes to train 20 epoches with AlexNet.

Please provide detailed descriptions of:

(1) which layers of AlexNet have been replaced

(2) the architecture of the new layers added including activation methods (same as problem 1)

(3) the final accuracy on test set along with time consumption for both training and testing

{Part 2: 20 points} Implement Strategy A where you use the activations of the pre-trained network as features to train one-vs-all SVMs for your scene classification task. Report the final accuracy on test set along with time consumption for both training and testing.

{Bonus: 10 points} Bonus will be given to those who fine-tune the [VGG network paper](#) and compare performance with AlexNet. Explain why VGG performed better or worse.

Hints:

- Many pre-trained models are available in PyTorch at [here](#).
- For fine-tuning pretrained network using PyTorch, please read this [tutorial](#).

```
1  # reload data with a larger size
2  img_size = (224, 224)
3  batch_num = 50 # training sample number per batch
4
5  # load training dataset
6  trainloader_large = list(load_dataset('/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw3/data/train/', img_size, batch_num=batch_num, shuffle=True, augment=True))
7  file = open('trainloader_large.pkl', 'wb')
8  pickle.dump(trainloader_large, file)
9  file.close()
10 train_num = len(trainloader_large)
11 print("Finish loading %d minibatches(=%d) of training samples." % (train_num, batch_num))
12
13 # load testing dataset
14 testloader_large = list(load_dataset('/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw3/data/test/', img_size, num_per_class=50, batch_num=batch_num, is_cifar10=True))
15 file = open('testloader_large.pkl', 'wb')
16 pickle.dump(testloader_large, file)
17 file.close()
18 test_num = len(testloader_large)
19 print("Finish loading %d minibatches(=%d) of testing samples." % (test_num, batch_num))
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```



```
AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=9216, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=4096, out_features=16, bias=True)
  )
)
```

```
1 def train(epoch, training_set, model):
2     for i in range(len(training_set)):
3         x_train, y_train = Variable(training_set[i][0]), Variable(training_set[i][1])
4         output_train = model(x_train.cuda(device))
5         loss_train = criterion(output_train, y_train.cuda(device))
6         optimizer.zero_grad()
7         loss_train.backward()
8         optimizer.step()
9
10 def test(training_set, model):
11     with torch.no_grad():
12         total = 0
13         correct = 0
14         for i in range(len(training_set)):
15             x_test = Variable(training_set[i][0])
16             outputs = model(x_test.cuda(device))
17             _, predicted = torch.max(outputs.data, 1)
18             total += training_set[i][1].size(0)
19             correct += (predicted == training_set[i][1].cuda(device)).sum().item()
20     print('Test Accuracy of the model on test images: {} %'.format((correct / total) * 100))
```

```
1 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
2 model = alexnet
3 model.cuda(device)
4 optimizer = Adam(model.parameters(), lr=0.0001)
5 criterion = CrossEntropyLoss()
```

```
6 start_time = time.time()
7 for epoch in range(20):
8     train(epoch, trainloader_large, model)
9 print("Time taken for training: ", round(time.time() - start_time), "seconds")
```

📄 Time taken for training: 119 seconds

```
1 start_time = time.time()
2 test(testloader_large, model)
3 print("Time taken for testing: ", round(time.time() - start_time), "seconds")
```

📄 Test Accuracy of the model on test images: 83.75 %
Time taken for testing: 0 seconds

(1) The fully connected - Linear layer (final) of AlexNet having **in_features = 4096** and **out_features = 1000**

```
Linear(in_features=4096, out_features=1000, bias=True)
```

was replaced.

(2) A new fully connected layer was added to give 16 outputs for each of our 16 classes.

```
Linear(in_features=4096, out_features=16, bias=True)
```

(3) The changed alexnet model has been trained and tested using cuda cores. The testing accuracy obtained = **83.75%** when the time taken for training was **119 seconds** and the time taken for testing was **~0 seconds**

Resources Referred:

<https://medium.com/@kushajreal/training-alexnet-with-tips-and-checks-on-how-to-train-cnns-practical-cnns-in-pytorch-1-61daa679c74a>

<https://neurohive.io/en/popular-networks/alexnet-imagenet-classification-with-deep-convolutional-neural-networks/>

```
1 alexnet = models.alexnet(pretrained=True)
2 alexnet.classifier = alexnet.classifier[:-2]
3 alexnet
```

📄

```
AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=9216, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
  )
)
```

```
1 def get_output(training_set, model):
2     output = []
3     for i in range(len(training_set)):
4         output_train = model(Variable(training_set[i][0]).cuda(device))
5         output.append(output_train)
6     return output
7
8 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
9 model = alexnet
10 model.cuda(device)
11 start_time = time.time()
12 alexnet_output = get_output(trainloader_large, model)
13 print("Time taken for training: ", round(time.time() - start_time), "seconds")
14 file = open('alexnet_output.pkl', 'wb')
15 pickle.dump(alexnet_output, file)
16 file.close()
```

 Time taken for training: 1 seconds

```
1 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
2 model = alexnet
3 model.cuda(device)
4 start_time = time.time()
5 alexnet_output_test = get_output(testloader_large, model)
6 file = open('alexnet_output_test.pkl', 'wb')
7 pickle.dump(alexnet_output_test, file)
8 file.close()
9 print("Time taken for extracting testing image descriptors: ", round(time.time() - start_time), "seconds")
```

```
➤ Time taken for extracting testing image descriptors:  0 seconds
```

```
1  file = open('alexnet_output.pkl', 'rb')
2  alexnet_output = pickle.load(file)
3  file.close()
4  file = open('alexnet_output_test.pkl', 'rb')
5  alexnet_output_test = pickle.load(file)
6  file.close()
7  alex_op_train = []
8  alex_op_train_labels = []
9  alex_op_test = []
10 alex_op_test_labels = []
11 for i in range(len(alexnet_output)):
12     for j in range(len(alexnet_output[i])):
13         alex_op_train.append(alexnet_output[i][j][:200].cpu().detach().numpy())
14         alex_op_train_labels.append(trainloader_large[i][1][j].item())
15 for i in range(len(alexnet_output_test)):
16     for j in range(len(alexnet_output_test[i])):
17         alex_op_test.append(alexnet_output_test[i][j][:200].cpu().detach().numpy())
18         alex_op_test_labels.append(testloader_large[i][1][j].item())
```

```
1  from sklearn.svm import LinearSVC
2  from sklearn.metrics import accuracy_score
3  linearSVCClassifier = LinearSVC(C=0.0001, multi_class="ovr")
4  linearSVCClassifier.fit(alex_op_train, alex_op_train_labels)
5  pred = linearSVCClassifier.predict(alex_op_test)
6  print('Accuracy: ', accuracy_score(pred, alex_op_test_labels)*100, '%')
```

```
➤ Accuracy:  80.25 %
```

After training the Linear Support Vector Classifier with the first 200 activations generated for each image in the penultimate layer in alexnet, we get a test accuracy of **80.25%** when the model is tested with the activations obtained for the test images.

```
1  vgg = models.vgg11(pretrained=True)
2  vgg.classifier[6] = Linear(in_features=4096, out_features=16, bias=True)
3  vgg
```

```
➤
```

Downloading: "<https://download.pytorch.org/models/vgg11-bbd30ac9.pth>" to /root/.cache/torch/checkpoints/vgg11-bbd30ac9.pth

100%|██████████| 507M/507M [00:08<00:00, 65.3MB/s]

```
VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (11): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (12): ReLU(inplace=True)
    (13): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (14): ReLU(inplace=True)
    (15): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (16): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (17): ReLU(inplace=True)
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (19): ReLU(inplace=True)
    (20): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
  (classifier): Sequential(
    (0): Linear(in_features=25088, out_features=4096, bias=True)
    (1): ReLU(inplace=True)
    (2): Dropout(p=0.5, inplace=False)
    (3): Linear(in_features=4096, out_features=4096, bias=True)
    (4): ReLU(inplace=True)
    (5): Dropout(p=0.5, inplace=False)
    (6): Linear(in_features=4096, out_features=16, bias=True)
  )
)
```

```
1 def trainVGG(epoch, training_set, model):
2     for i in range(len(training_set)):
3         x_train, y_train = Variable(training_set[i][0]), Variable(training_set[i][1])
4         output_train = model(x_train.cuda(device))
5         loss_train = criterion(output_train, y_train.cuda(device))
6         optimizer.zero_grad()
7         loss_train.backward()
8         optimizer.step()
9
10 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
11 model = vgg
12 model.cuda(device)
13 optimizer = Adam(model.parameters(), lr=0.0001)
14 criterion = CrossEntropyLoss()
15 start_time = time.time()
16 for epoch in range(20):
17     trainVGG(epoch, trainloader_large, model)
18 print("Time taken for training: ". round(time.time() - start time), "seconds")
```

```
Time taken for training: 740 seconds
```

```
1 def testVGG(training_set, model):
2     with torch.no_grad():
3         total = 0
4         correct = 0
5         for i in range(len(training_set)):
6             x_test = Variable(training_set[i][0])
7             outputs = model(x_test.cuda(device))
8             _, predicted = torch.max(outputs.data, 1)
9             total += training_set[i][1].size(0)
10            correct += (predicted == training_set[i][1].cuda(device)).sum().item()
11            print('Test Accuracy of the model on test images: {} %'.format((correct / total) * 100))
12 start_time = time.time()
13 testVGG(testloader_large, model)
14 print("Time taken for testing: ", round(time.time() - start_time), "seconds")
```

```
Test Accuracy of the model on test images: 87.25 %
Time taken for testing: 2 seconds
```

A VGG net was trained and tested with the same data that was passed to the alexnet on cuda cores.

The VGG Net's last layer was replaced in the similar way as done before.

Time taken for training VGG net = **740 seconds**

Though the training time on the VGG net was significantly larger as compared to AlexNet, we observe a substantial increase in the testing accuracy of the model.

VGG Net gives us an accuracy of **87.25%** which is an improvement of **3.5%** over AlexNet.

In VGG Net, the convolution kernels are of size 3x3 as compared to multiple kernel sizes in AlexNet (11x11, 5x5, 3x3).

These variable size kernels can be replicated by the 3x3 kernels in VGG which leads to reduction in the number of trainable variables. This property makes VGG Net more robust to over-fitting, which might be the main cause behind the increase in testing accuracy.

Moreover, the number of strides in the first convolutional layer of AlexNet is 4. This causes information loss in training of AlexNet which might result in a reduced testing accuracy.

Resource Referred: <https://towardsdatascience.com/the-w3h-of-alexnet-vggnet-resnet-and-inception-7baaecccc96>

Submission guidelines

Extract the downloaded .zip file to a folder of your preference. The input and output paths are predefined and **DO NOT** change them, (we assume that 'Surname_Givenname_SBUID_hw4' is your working directory, and all the paths are relative to this directory). The image read and write functions are already written for you. All you need to do is to fill in the blanks as indicated to generate proper outputs. **DO NOT** zip and upload the dataset on blackboard due to size limit.

When submitting your .zip file through blackboard, please -- name your .zip file as **Surname_Givenname_SBUID_hw*.zip**.

This zip file should include:

```
Surname_Givenname_SBUID_hw*
|---Surname_Givenname_SBUID_hw*.ipynb
|---Surname_Givenname_SBUID_hw*.py
|---Surname_Givenname_SBUID_hw*.pdf
```

where Surname_Givenname_SBUID_hw.py is the Python code of Surname_Givenname_SBUID_hw.ipynb, which can be dowloaded by File->Download .py.

For instance, student Michael Jordan should submit a zip file named "Jordan_Michael_111134567_hw4.zip" for homework4 in this structure:

```
Jordan_Michael_111134567_hw4
|---Jordan_Michael_111134567_hw4.ipynb
|---Jordan_Michael_111134567_hw4.py
|---Jordan_Michael_111134567_hw4.pdf
```

The **Surname_Givenname_SBUID_hw*.pdf** should include a **google shared link** and **Surname_Givenname_SBUID_Pred*.pdf** should be your test set prediction file in the specified format. To generate the **google shared link**, first create a folder named **Surname_Givenname_SBUID_hw*** in your Google Drive with your Stony Brook account. The structure of the files in the folder should be exactly the same as the one you downloaded. If you alter the folder structures, the grading of your homework will be significantly delayed and possibly penalized.

Then right click this folder, click **Get shareable link**, in the People textfield, enter two TA's emails: bo.cao.1@stonybrook.edu and sayontan.ghosh@stonybrook.edu. Make sure that TAs who have the link **can edit, not just can view**, and also **uncheck** the **Notify people** box.

Colab has a good feature of version control, you should take advantage of this to save your work properly. However, the timestamp of the submission made in blackboard is the only one that we consider for grading. To be more specific, we will only grade the version of your code right before the timestamp of the submission made in blackboard.

You are encouraged to post and answer questions on Piazza. Based on the amount of email that we have received in past years, there might be dealys in replying to personal emails. Please ask questions on Piazza and send emails only for personal issues.

Be aware that your code will undergo plagiarism check both vertically and horizontally. Please do your own work.

Late submission penalty:

There will be a 10% penalty per day for late submission. However, you will have 4 days throughout the whole semester to submit late without penalty. Note that the grace period is calculated by days instead of hours. If you submit the homework one minute after the deadline, one late day will be counted. Likewise, if you submit one minute after the deadline, the 10% penalty will be imposed if not using the grace period.

Drive Link: <https://drive.google.com/drive/folders/1Z1J0tPbXFxJVgAom7Xpnuk9DVjHxYLjB?usp=sharing>

