

```
1 from google.colab import drive
2 drive.mount('/content/gdrive', force_remount=True)
```

➞ Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.ap

Enter your authorization code:

.....

Mounted at /content/gdrive

```
1 %pwd
```

➞ '/content'

```
1 cd '/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw5'
```

➞ /content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw5

```
1 %pwd
```

```
2 %ls -lrt
```

➞

```
total 8455342
-rw----- 1 root root 8658247680 Oct 18 17:35 UCF101_images.tar
drwx----- 2 root root      4096 Nov 20 11:13 annos/
-rw----- 1 root root      17664 Nov 21 03:49 CSE527_HW5_fall19.ipynb
```

```
1 !wget 'http://vision.cs.stonybrook.edu/~yangwang/public/UCF101_images.tar'
```

➞

```
--2019-11-21 03:39:18--  http://vision.cs.stonybrook.edu/~yangwang/public/UCF101_images.tar
Resolving vision.cs.stonybrook.edu (vision.cs.stonybrook.edu)... 130.245.4.232
Connecting to vision.cs.stonybrook.edu (vision.cs.stonybrook.edu)|130.245.4.232|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 8658247680 (8.1G) [application/x-tar]
Saving to: 'UCF101_images.tar.1'
```

```
UCF101_images.tar.1 100%[=====>]    8.06G  31.9MB/s    in 4m 30s
```

```
2019-11-21 03:43:48 (30.6 MB/s) - 'UCF101_images.tar.1' saved [8658247680/8658247680]
```

```

1  !tar -xkf './UCF101_images.tar' 2>/dev/null

1  cd '/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw5/'

📁 /content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw5

1  %ls -lrt

📁 total 8455446
drwx----- 13322 root root      4096 May 31  2017 images/
-rw-----   1 root root 8658247680 Oct 18 17:35 UCF101_images.tar
drwx-----   2 root root      4096 Nov 20 11:13 annos/
-rw-----   1 root root    66439 Nov 21 08:20 train_video_label_df.pkl
-rw-----   1 root root    26417 Nov 21 08:20 test_video_label_df.pkl
-rw-----   1 root root    26921 Nov 21 08:31 Paratkar_Shreyash_112673930_hw5.ipynb

```

▼ Action Recognition @ UCF101

Due date: 11:59 pm on Nov. 19, 2019 (Tuesday)

Description

In this homework, you will be doing action recognition using Recurrent Neural Network (RNN), (Long-Short Term Memory) LSTM in particular. You will be given a dataset called UCF101, which consists of 101 different actions/classes and for each action, there will be 145 samples. We tagged each sample into either training or testing. Each sample is supposed to be a short video, but we sampled 25 frames from each videos to reduce the amount of data. Consequently, a training sample is an image tuple that forms a 3D volume with one dimension encoding *temporal correlation* between frames and a label indicating what action it is.

To tackle this problem, we aim to build a neural network that can not only capture spatial information of each frame but also temporal information between frames. Fortunately, you don't have to do this on your own. RNN — a type of neural network designed to deal with time-series data — is right here for you to use. In particular, you will be using LSTM for this task.

Instead of training an end-to-end neural network from scratch whose computation is prohibitively expensive, we divide this into two steps: feature extraction and modelling. Below are the things you need to implement for this homework:

- **{35 pts} Feature extraction.** Use any of the [pre-trained models](#) to extract features from each frame. Specifically, we recommend not to use the activations of the last layer as the features tend to be task specific towards the end of the network. **hints:**
 - A good starting point would be to use a pre-trained VGG16 network, we suggest first fully connected layer `torchvision.models.vgg16` (4096 dim) as features of each video frame. This will result into a 4096x25 matrix for each video.
 - Normalize your images using `torchvision.transforms`

```
normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
prep = transforms.Compose([ transforms.ToTensor(), normalize ])
prep(img)
```

The mean and std. mentioned above is specific to Imagenet data

More details of image preprocessing in PyTorch can be found at http://pytorch.org/tutorials/beginner/data_loading_tutorial.html

- **{35 pts} Modelling.** With the extracted features, build an LSTM network which takes a **dx25** sample as input (where **d** is the dimension of the extracted feature for each frame), and outputs the action label of that sample.
- **{20 pts} Evaluation.** After training your network, you need to evaluate your model with the testing data by computing the prediction accuracy (**5 points**). The baseline test accuracy for this data is 75%, and **10 points** out of 20 is for achieving test accuracy greater than the baseline. Moreover, you need to compare (**5 points**) the result of your network with that of support vector machine (SVM) (stacking the **dx25** feature matrix to a long vector and train a SVM).
- **{10 pts} Report.** Details regarding the report can be found in the submission section below.

Notice that the size of the raw images is 256x340, whereas your pre-trained model might take **nxn** images as inputs. To solve this problem, instead of resizing the images which unfavorably changes the spatial ratio, we take a better solution: Cropping five **nxn** images, one at the image center and four at the corners and compute the **d**-dim features for each of them, and average these five **d**-dim feature to get a final feature representation for the raw image. For example, VGG takes 224x224 images as inputs, so we take the five 224x224 croppings of the image, compute 4096-dim VGG features for each of them, and then take the mean of these five 4096-dim vectors to be the representation of the image.

In order to save you computational time, you need to do the classification task only for **the first 25** classes of the whole dataset. The same applies to those who have access to GPUs. **Bonus 10 points for running and reporting on the entire 101 classes.**

Dataset

Download **dataset** at [UCF101](#)(Image data for each video) and the **annos folder** which has the video labels and the label to class name mapping is included in the assignment folder uploaded.

UCF101 dataset contains 101 actions and 13,320 videos in total.

- annos/actions.txt
 - lists all the actions (ApplyEyeMakeup , .., YoYo)
- annots/videos_labels_subsets.txt
 - lists all the videos (v_000001 , .., v_013320)
 - labels (1 , .., 101)
 - subsets (1 for train, 2 for test)
- images/
 - each folder represents a video
 - the video/folder name to class mapping can be found using annots/videos_labels_subsets.txt , for e.g. v_000001 belongs to class 1 i.e. ApplyEyeMakeup
 - each video folder contains 25 frames

Some Tutorials

- Good materials for understanding RNN and LSTM
 - <http://blog.echen.me>
 - <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>
 - <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- Implementing RNN and LSTM with PyTorch
 - [LSTM with PyTorch](#)
 - [RNN with PyTorch](#)

▼ Problem 1. Feature extraction

```
1  # \*write your codes for feature extraction (You can use multiple cells, this is just a place holder)
2  import os
3  os.environ['CUDA_LAUNCH_BLOCKING'] = '1'
4  import torchvision.models as models
5  from __future__ import print_function, division
6  import torch
7  import pandas as pd
8  from skimage import io, transform
9  import numpy as np
10 import matplotlib.pyplot as plt
11 from torch.utils.data import Dataset, DataLoader
12 from torchvision import transforms, utils
13 import pickle
14 import cv2
15 from torch.autograd import Variable
16 import torch.nn as nn
17 import torch.nn.functional as F
18 import torch.optim as optim
19 # Ignore warnings
20 import warnings
21 import time
22 warnings.filterwarnings("ignore")
23
24 plt.ion()    # interactive mode

1  vgg16 = models.vgg16(pretrained=True)
2  vgg16.classifier = vgg16.classifier[:2]
3  vgg16
```



Downloading: "<https://download.pytorch.org/models/vgg16-397923af.pth>" to /root/.cache/torch/checkpoints/vgg16-397923af.pth
100%|██████████| 528M/528M [00:21<00:00, 25.8MB/s]

```
VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace=True)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace=True)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace=True)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace=True)
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): ReLU(inplace=True)
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (18): ReLU(inplace=True)
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (20): ReLU(inplace=True)
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (22): ReLU(inplace=True)
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (25): ReLU(inplace=True)
    (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (27): ReLU(inplace=True)
    (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (29): ReLU(inplace=True)
    (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
  (classifier): Sequential(
    (0): Linear(in_features=25088, out_features=4096, bias=True)
    (1): ReLU(inplace=True)
  )
)
```

```
1 videos_labels_subsets = open("annos/videos_labels_subsets.txt". "r")
```

```

1 videos_labels_subsets = open('annos/videos_labels_subsets.txt', 'r')
2 dic = {}
3 train_list = []
4 test_list = []
5 cnt=0
6 for l in videos_labels_subsets:
7     lst = l[:-1].split('\t')
8     if int(lst[1]) < 21 or int(lst[1]) > 25:
9         continue
10    dic['v_num'] = lst[0]
11    dic['class'] = lst[1]
12    if lst[2]=='1':
13        train_list.append(dic)
14    else:
15        test_list.append(dic)
16    dic = {}
17 train_video_label_df = pd.DataFrame(train_list)
18 test_video_label_df = pd.DataFrame(test_list)
19 file = open('train_video_label_df.pkl','wb')
20 pickle.dump(train_video_label_df, file)
21 file.close()
22 file = open('test_video_label_df.pkl','wb')
23 pickle.dump(test_video_label_df, file)
24 file.close()
25 print(len(train_list))
26 print(len(test_list))

```



476
190

```

1 file = open('train_video_label_df.pkl', 'rb')
2 train_video_label_df = pickle.load(file)
3 file.close()
4 file = open('test_video_label_df.pkl', 'rb')
5 test_video_label_df = pickle.load(file)
6 file.close()

```

```

1 import os
2 videos_labels_subsets2 = open("annos/videos_labels_subsets.txt", "r")
3 img_count = 0
4 for l in videos_labels_subsets2:

```

```

5     lst = l[:-1].split('\t')
6     if int(lst[1]) > 25:
7         break
8     path, dirs, files = next(os.walk('/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_1126739
9     img_count = img_count + len(files)
10    img_count

```

```

1    root_dir = '/content/gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/Paratkar_Shreyash_112673930_hw5/images/'
2    # image = cv2.imread(root_dir+'v_000001/i_0001.jpg')
3    # # print(torch.tensor(image).shape)
4    normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
5    prep = transforms.Compose([ transforms.ToTensor(), normalize ])
6    # image = prep(image)
7    # print(image.shape)
8    # img_tp_lt = image[0:3, :224, :224]
9    # print(img_tp_lt.shape)
10   # img_bt_lt = image[0:3, 32:, :224]
11   # print(img_bt_lt.shape)
12   # img_bt_rt = image[0:3, 32:, 116:]
13   # print(img_bt_rt.shape)
14   # img_tp_rt = image[0:3, :224, 116:]
15   # print(img_tp_rt.shape)
16   # img_cn = image[0:3, 16:240, 58:282]
17   # print(img_cn.shape)
18   # labels = [1] * 125
19   # labels = torch.Tensor(np.array(labels))
20   # print(labels)

```

```

1    print(torch.cuda.memory_cached()-torch.cuda.memory_allocated())

```

0

```

1    class UCF101ImageDataset(Dataset):
2        """UCF101 Actions dataset."""
3
4        def __init__(self, pickle_file, root_dir, transform=transforms.Compose([
5            transforms.ToTensor(),
6            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
7            ])):
8        """

```


Args:

csv_file (string): Path to the csv file with annotations.

root_dir (string): Directory with all the images.

transform (callable, optional): Optional transform to be applied
on a sample.

"""

file = open(pickle_file, 'rb')

self.actions_frame = pickle.load(file)

self.root_dir = root_dir

self.transform = transform

def __len__(self):

return len(self.actions_frame)

def __getitem__(self, idx):

action_num = self.actions_frame.iloc[idx, 1]

image_transforms = []

for i in range(1, 26):

img_name = ''

if i<10:

img_name = 'i_000'+str(i)+'.jpg'

else:

img_name = 'i_00'+str(i)+'.jpg'

parent_img_name = self.root_dir+self.actions_frame.iloc[idx, 0]+ '/' + img_name

image = cv2.imread(parent_img_name)

print(parent_img_name)

image = prep(image)

print(image.shape)

image_transforms.append(image[0:3, :224, :224])

print(img_tp_lt.shape)

image_transforms.append(image[0:3, 32:, :224])

print(img_bt_lt.shape)

image_transforms.append(image[0:3, 32:, 116:])

print(img_bt_rt.shape)

image_transforms.append(image[0:3, :224, 116:])

print(img_tp_rt.shape)

image_transforms.append(image[0:3, 16:240, 58:282])

print(img_cn.shape)

label = int(action_num)

print(len(image_transforms))

print(torch.stack(image_transforms).shape)

```

47         # print(torch.stack(image_transforms).shape)
50         sample = {'images':torch.stack(image_transforms), 'action':label}
51         return sample
52 train_dataset_func = UCF101ImageDataset(pickle_file='train_video_label_df.pkl', root_dir=root_dir)
53 test_dataset_func = UCF101ImageDataset(pickle_file='test_video_label_df.pkl', root_dir=root_dir)
54 train_dataloader = DataLoader(train_dataset_func, batch_size=1, shuffle=False, num_workers=4)
55 test_dataloader = DataLoader(test_dataset_func, batch_size=1, shuffle=False, num_workers=4)

```

```

1 # torch.cuda.empty_cache()
2 !nvidia-smi

```

 Fri Nov 22 07:59:51 2019

```

+-----+
| NVIDIA-SMI 430.50      Driver Version: 418.67      CUDA Version: 10.1      |
+-----+-----+-----+-----+-----+
| GPU   Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
+-----+-----+-----+-----+-----+
|    0   Tesla K80           Off   | 00000000:00:04.0 Off |                    0 |
| N/A   35C    P8      26W / 149W | 11MiB / 11441MiB |      0%      Default |
+-----+-----+-----+-----+-----+

```

```

+-----+
| Processes:                                     GPU Memory |
|  GPU       PID    Type    Process name                     Usage      |
+-----+-----+-----+-----+-----+
| No running processes found                       |
+-----+

```

```

1 import time
2 model = vgg16
3 model = model.cuda()
4 count = 0
5 train_features_data = []
6 labels = []
7 model.eval()
8 start_time = time.time()
9 with torch.no_grad():
10     for video_num, sample in enumerate(test_dataloader):
11         print("Video taken: ", video_num)
12         print("Time since start ", round(time.time() - start_time), " seconds")
13         video_images_tensor = sample['images'][0].cuda()
14         label = sample['action'][0]

```

```
14     label = sample['action'][0].cuda()
15
16     vgg16output_125 = model(video_images_tensor)
17
18     del video_images_tensor
19
20     for i in range(0, 125, 5):
21         tensor = vgg16output_125[i]
22         # print("First tensor: ", tensor)
23         for j in range(1,5):
24             tensor+=vgg16output_125[i+j]
25         # print("Mean tensor: ", tensor)
26         train_features_data.append(tensor/5)
27         labels.append(label)
28         del tensor
29     del vgg16output_125, label
30     print("Length of training features: ", len(train_features_data))
31     print("Length labels: ", len(labels))
```



Video taken: 0
Time since start 22 seconds
Video taken: 1
Time since start 24 seconds
Video taken: 2
Time since start 25 seconds
Video taken: 3
Time since start 26 seconds
Video taken: 4
Time since start 28 seconds
Video taken: 5
Time since start 29 seconds
Video taken: 6
Time since start 31 seconds
Video taken: 7
Time since start 32 seconds
Video taken: 8
Time since start 34 seconds
Video taken: 9
Time since start 36 seconds
Video taken: 10
Time since start 37 seconds
Video taken: 11
Time since start 38 seconds
Video taken: 12
Time since start 40 seconds
Video taken: 13
Time since start 41 seconds
Video taken: 14
Time since start 43 seconds
Video taken: 15
Time since start 44 seconds
Video taken: 16
Time since start 46 seconds
Video taken: 17
Time since start 47 seconds
Video taken: 18
Time since start 49 seconds
Video taken: 19
Time since start 50 seconds
Video taken: 20
Time since start 53 seconds
Video taken: 21
Time since start 54 seconds
Video taken: 22
Time since start 56 seconds

Time since start	56	seconds
Video taken:	23	
Time since start	57	seconds
Video taken:	24	
Time since start	59	seconds
Video taken:	25	
Time since start	61	seconds
Video taken:	26	
Time since start	62	seconds
Video taken:	27	
Time since start	64	seconds
Video taken:	28	
Time since start	65	seconds
Video taken:	29	
Time since start	67	seconds
Video taken:	30	
Time since start	68	seconds
Video taken:	31	
Time since start	70	seconds
Video taken:	32	
Time since start	71	seconds
Video taken:	33	
Time since start	73	seconds
Video taken:	34	
Time since start	75	seconds
Video taken:	35	
Time since start	76	seconds
Video taken:	36	
Time since start	79	seconds
Video taken:	37	
Time since start	80	seconds
Video taken:	38	
Time since start	82	seconds
Video taken:	39	
Time since start	83	seconds
Video taken:	40	
Time since start	85	seconds
Video taken:	41	
Time since start	87	seconds
Video taken:	42	
Time since start	88	seconds
Video taken:	43	
Time since start	89	seconds
Video taken:	44	
Time since start	91	seconds
Video taken:	45	

Time since start	95	seconds
Video taken:	46	
Time since start	97	seconds
Video taken:	47	
Time since start	98	seconds
Video taken:	48	
Time since start	99	seconds
Video taken:	49	
Time since start	101	seconds
Video taken:	50	
Time since start	103	seconds
Video taken:	51	
Time since start	104	seconds
Video taken:	52	
Time since start	105	seconds
Video taken:	53	
Time since start	108	seconds
Video taken:	54	
Time since start	109	seconds
Video taken:	55	
Time since start	110	seconds
Video taken:	56	
Time since start	111	seconds
Video taken:	57	
Time since start	113	seconds
Video taken:	58	
Time since start	115	seconds
Video taken:	59	
Time since start	116	seconds
Video taken:	60	
Time since start	117	seconds
Video taken:	61	
Time since start	120	seconds
Video taken:	62	
Time since start	121	seconds
Video taken:	63	
Time since start	122	seconds
Video taken:	64	
Time since start	123	seconds
Video taken:	65	
Time since start	126	seconds
Video taken:	66	
Time since start	127	seconds
Video taken:	67	
Time since start	128	seconds
Video taken:	68	

Time since start	130	seconds
Video taken:	69	
Time since start	132	seconds
Video taken:	70	
Time since start	133	seconds
Video taken:	71	
Time since start	135	seconds
Video taken:	72	
Time since start	136	seconds
Video taken:	73	
Time since start	138	seconds
Video taken:	74	
Time since start	140	seconds
Video taken:	75	
Time since start	141	seconds
Video taken:	76	
Time since start	142	seconds
Video taken:	77	
Time since start	145	seconds
Video taken:	78	
Time since start	146	seconds
Video taken:	79	
Time since start	147	seconds
Video taken:	80	
Time since start	148	seconds
Video taken:	81	
Time since start	152	seconds
Video taken:	82	
Time since start	153	seconds
Video taken:	83	
Time since start	154	seconds
Video taken:	84	
Time since start	155	seconds
Video taken:	85	
Time since start	158	seconds
Video taken:	86	
Time since start	159	seconds
Video taken:	87	
Time since start	160	seconds
Video taken:	88	
Time since start	161	seconds
Video taken:	89	
Time since start	164	seconds
Video taken:	90	
Time since start	165	seconds
Video taken:	91	

Time since start 166 seconds
Video taken: 92
Time since start 167 seconds
Video taken: 93
Time since start 169 seconds
Video taken: 94
Time since start 171 seconds
Video taken: 95
Time since start 172 seconds
Video taken: 96
Time since start 173 seconds
Video taken: 97
Time since start 175 seconds
Video taken: 98
Time since start 177 seconds
Video taken: 99
Time since start 178 seconds
Video taken: 100
Time since start 179 seconds
Video taken: 101
Time since start 182 seconds
Video taken: 102
Time since start 183 seconds
Video taken: 103
Time since start 184 seconds
Video taken: 104
Time since start 185 seconds
Video taken: 105
Time since start 188 seconds
Video taken: 106
Time since start 189 seconds
Video taken: 107
Time since start 190 seconds
Video taken: 108
Time since start 191 seconds
Video taken: 109
Time since start 194 seconds
Video taken: 110
Time since start 195 seconds
Video taken: 111
Time since start 196 seconds
Video taken: 112
Time since start 198 seconds
Video taken: 113
Time since start 201 seconds
Video taken: 114

video taken: 114
Time since start 202 seconds
Video taken: 115
Time since start 203 seconds
Video taken: 116
Time since start 204 seconds
Video taken: 117
Time since start 207 seconds
Video taken: 118
Time since start 208 seconds
Video taken: 119
Time since start 209 seconds
Video taken: 120
Time since start 210 seconds
Video taken: 121
Time since start 213 seconds
Video taken: 122
Time since start 214 seconds
Video taken: 123
Time since start 215 seconds
Video taken: 124
Time since start 216 seconds
Video taken: 125
Time since start 219 seconds
Video taken: 126
Time since start 220 seconds
Video taken: 127
Time since start 221 seconds
Video taken: 128
Time since start 222 seconds
Video taken: 129
Time since start 224 seconds
Video taken: 130
Time since start 226 seconds
Video taken: 131
Time since start 227 seconds
Video taken: 132
Time since start 228 seconds
Video taken: 133
Time since start 230 seconds
Video taken: 134
Time since start 232 seconds
Video taken: 135
Time since start 233 seconds
Video taken: 136
Time since start 234 seconds

Video taken: 137
Time since start 237 seconds
Video taken: 138
Time since start 238 seconds
Video taken: 139
Time since start 239 seconds
Video taken: 140
Time since start 240 seconds
Video taken: 141
Time since start 243 seconds
Video taken: 142
Time since start 244 seconds
Video taken: 143
Time since start 245 seconds
Video taken: 144
Time since start 246 seconds
Video taken: 145
Time since start 250 seconds
Video taken: 146
Time since start 251 seconds
Video taken: 147
Time since start 252 seconds
Video taken: 148
Time since start 253 seconds
Video taken: 149
Time since start 256 seconds
Video taken: 150
Time since start 257 seconds
Video taken: 151
Time since start 258 seconds
Video taken: 152
Time since start 259 seconds
Video taken: 153
Time since start 263 seconds
Video taken: 154
Time since start 265 seconds
Video taken: 155
Time since start 266 seconds
Video taken: 156
Time since start 267 seconds
Video taken: 157
Time since start 270 seconds
Video taken: 158
Time since start 271 seconds
Video taken: 159
Time since start 272 seconds

Video taken: 160
Time since start 273 seconds
Video taken: 161
Time since start 276 seconds
Video taken: 162
Time since start 277 seconds
Video taken: 163
Time since start 279 seconds
Video taken: 164
Time since start 280 seconds
Video taken: 165
Time since start 285 seconds
Video taken: 166
Time since start 287 seconds
Video taken: 167
Time since start 288 seconds
Video taken: 168
Time since start 289 seconds
Video taken: 169
Time since start 292 seconds
Video taken: 170
Time since start 293 seconds
Video taken: 171
Time since start 294 seconds
Video taken: 172
Time since start 295 seconds
Video taken: 173
Time since start 298 seconds
Video taken: 174
Time since start 299 seconds
Video taken: 175
Time since start 301 seconds
Video taken: 176
Time since start 302 seconds
Video taken: 177
Time since start 305 seconds
Video taken: 178
Time since start 306 seconds
Video taken: 179
Time since start 307 seconds
Video taken: 180
Time since start 308 seconds
Video taken: 181
Time since start 312 seconds
Video taken: 182
Time since start 313 seconds

```
Video taken: 183
Time since start 314 seconds
Video taken: 184
Time since start 316 seconds
Video taken: 185
Time since start 319 seconds
Video taken: 186
Time since start 320 seconds
Video taken: 187
Time since start 321 seconds
Video taken: 188
Time since start 322 seconds
Video taken: 189
Time since start 326 seconds
Length of training features: 4750
Length labels: 4750
```

```
1 # file = open('test_features_data_class_21_25.pkl','wb')
2 # pickle.dump(train_features_data, file)
3 # file.close()
4 # file = open('test_labels_class_21_25.pkl','wb')
5 # pickle.dump(labels, file)
6 # file.close()
7 print("Length of training features: ", len(train_features_data))
8 print("Length labels: ", len(labels))
9 # file = open('train_features_data_class_21_25.pkl','wb')
10 # pickle.dump(train_features_data, file)
11 # file.close()
12 # file = open('labels_class_21_25.pkl','wb')
13 # pickle.dump(labels, file)
14 file.close()
```

☞ Length of training features: 4750
Length labels: 4750

```
1 # file = open('test_features_data_class_11_15.pkl', 'rb')
2 # test_features_data_class_11_15 = pickle.load(file)
3 # print(len(test_features_data_class_11_15))
4 # file.close()
5 # file = open('test_labels_class_11_15.pkl', 'rb')
6 # test_labels_class_11_15 = pickle.load(file)
```

```

7 # print(len(test_labels_class_11_15))
8 # file.close()
9 # file = open('train_features_data_class_1_5.pkl', 'rb')
10 # train_features_data_class_0_5 = pickle.load(file)
11 # print(len(train_features_data_class_0_5))
12 # file.close()
13 # file = open('train_features_data_class_1_5.pkl', 'rb')
14 # labels_class_0_5 = pickle.load(file)
15 # print(len(labels_class_0_5))
16 # file.close()

```



```

11525
11525

```

```

1 train_pkls = ['train_features_data_class_1_5.pkl', 'train_features_data_class_6_10.pkl', 'train_features_data_class_11_15.pkl', '
2 train_label_pkls = ['labels_class_1_5.pkl', 'labels_class_6_10.pkl', 'labels_class_11_15.pkl', 'labels_class_16_20.pkl', 'labels_
3 test_pkls = ['test_features_data_class_1_5.pkl', 'test_features_data_class_6_10.pkl', 'test_features_data_class_11_15.pkl', 'tes
4 test_label_pkls = ['test_labels_class_1_5.pkl', 'test_labels_class_6_10.pkl', 'test_labels_class_11_15.pkl', 'test_labels_class_1
5 train_list = []
6 train_labels = []
7 test_list = []
8 test_labels = []
9 for i in range(0,5):
10     train_list.extend(pickle.load(open(train_pkls[i], 'rb')))
11     train_labels.extend(pickle.load(open(train_label_pkls[i], 'rb')))
12     test_list.extend(pickle.load(open(test_pkls[i], 'rb')))
13     test_labels.extend(pickle.load(open(test_label_pkls[i], 'rb')))
14 print(len(train_list))
15 print(len(train_labels))
16 print(len(test_list))
17 print(len(test_labels))

```



```

60225
60225
23775
23775

```

```

1 train_data_final = torch.stack(train_list)
2 print("Training data size:\t", train_data_final.size(), "\t type: ", type(train_data_final), "\t\t element: ", (train_data_final[
3 train_labels_final = torch.stack(train_labels)
4 print("Training labels size:\t", train_labels_final.size(), "\t\t type: ", type(train_labels_final), "\t\t element: ", (train_labels_final[

```

```

4 print(' Training labels size:\t', train_labels_final.size(), '\t\t type: ', type(train_labels_final), '\t\t element: ', (train_labels_final[0]))
5 test_data_final = torch.stack(test_list)
6 print("Testing data size:\t", test_data_final.size(), "\t type: ", type(test_data_final), "\t\t element: ", (test_data_final[0]))
7 test_labels_final = torch.stack(test_labels)
8 print("Testing labels size:\t", test_labels_final.size(), "\t\t type: ", type(test_labels_final), "\t\t element: ", (test_labels_final[0]))

```

☞ Training data size:	torch.Size([60225, 4096])	type:	<class 'torch.Tensor'>	element:	tensor([0.0542, 1.0759, ..., 0.0000, 0.0000])
Training labels size:	torch.Size([60225])	type:	<class 'torch.Tensor'>	element:	tensor(1, device='cuda:0')
Testing data size:	torch.Size([23775, 4096])	type:	<class 'torch.Tensor'>	element:	tensor([0.0694, 0.0000, ..., 0.0000, 0.0000])
Testing labels size:	torch.Size([23775])	type:	<class 'torch.Tensor'>	element:	tensor(1, device='cuda:0')

Problem 2. Modelling

Print the size of your training and test data

```

1 # Don't hardcode the shape of train and test data
2 print('Shape of training data is :', )
3 print("Training data size:\t", train_data_final.size(), "\t shape: ", train_data_final.shape, "\t type: ", type(train_data_final))
4 print("Training labels size:\t", train_labels_final.size(), "\t\t shape: ", train_labels_final.shape, "\t\t type: ", type(train_labels_final))
5 print('Shape of test/validation data is :', )
6 print("Testing data size:\t", test_data_final.size(), "\t shape: ", test_data_final.shape, "\t type: ", type(test_data_final), "\t\t element: ", (test_data_final[0]))
7 print("Testing labels size:\t", test_labels_final.size(), "\t\t shape: ", test_labels_final.shape, "\t\t type: ", type(test_labels_final))

```

☞ Shape of training data is :			
Training data size:	torch.Size([60225, 4096])	shape:	torch.Size([60225, 4096])
Training labels size:	torch.Size([60225])	shape:	torch.Size([60225])
Shape of test/validation data is :			
Testing data size:	torch.Size([23775, 4096])	shape:	torch.Size([23775, 4096])
Testing labels size:	torch.Size([23775])	shape:	torch.Size([23775])

```

1 file = open('train_data_final.pkl','wb')
2 pickle.dump(train_data_final, file)
3 file.close()
4 file = open('train_labels_final.pkl','wb')
5 pickle.dump(train_labels_final, file)
6 file.close()
7 file = open('test_data_final.pkl','wb')
8 pickle.dump(test_data_final, file)

```



```

17 test_labels_final_split = torch.tensor(test_lab)
18 print(test_labels_final_split.shape)
19 del train_lab, train_data_final, train_labels_final, test_lab, test_data_final, test_labels_final

```

```

☞ torch.Size([2409, 25, 4096])
   torch.Size([2409])
   torch.Size([951, 25, 4096])
   torch.Size([951])

```

```

1  import random
2  class LSTM(nn.Module):
3
4      def __init__(self, input_dim, hidden_dim, batch_size, output_dim=1, num_layers=2):
5          super(LSTM, self).__init__()
6          self.input_dim = input_dim
7          self.hidden_dim = hidden_dim
8          self.batch_size = batch_size
9          self.num_layers = num_layers
10         self.lstm = nn.LSTM(self.input_dim, self.hidden_dim, self.num_layers)
11         self.linear = nn.Linear(self.hidden_dim, output_dim)
12
13     def init_hidden(self):
14         return (torch.zeros(self.num_layers, self.batch_size, self.hidden_dim), torch.zeros(self.num_layers, self.batch_size, self.hidden_dim))
15
16     def forward(self, input):
17         lstm_out, self.hidden = self.lstm(input.view(len(input), self.batch_size, -1))
18         y_pred = self.linear(lstm_out[-1].view(self.batch_size, -1))
19         return y_pred.view(-1)
20
21 def trainLSTM(model, video_features, video_labels, epochs, learning_rate, optimizer=None):
22     device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
23     optimizer = optim.Adam(model.parameters(), lr=learning_rate)
24     model = model.cuda(device)
25     loss_func = nn.CrossEntropyLoss()
26     start_time = time.time()
27     for ep in range(epochs):
28         data = list(zip(video_features, video_labels))
29         random.shuffle(data)
30         video_features, video_labels = zip(*data)
31         correct_count = 0
32         total_count = 0

```



```

33     for i in range(len(video_features)):
34         video_feature = video_features[i]
35         label_feature = video_labels[i].type(torch.LongTensor)
36         label_feature = label_feature.view(1) - 1
37         model.zero_grad()
38         model.hidden = model.init_hidden()
39         preds = model(video_feature.cuda(device))
40         preds = preds.view(1, preds.shape[0])
41         loss = loss_func(preds, label_feature.cuda(device))
42         loss.backward()
43         optimizer.step()
44         pred_indices, pred_vals = torch.max(preds.data, 1)
45         total_count = total_count + label_feature.size(0)
46         correct_count = correct_count + (pred_vals == label_feature.cuda(device)).sum().item()
47     torch.cuda.empty_cache()
48     print('Epoch: ', ep + 1, ' = ', round((correct_count / total_count) * 100, 3), '%')
49
50     print('\nTraining time = ', round(time.time() - start_time), ' seconds')
51     return (correct_count / total_count) * 100
52 model = LSTM(input_dim=4096, hidden_dim=1024, num_layers=2, batch_size=1, output_dim = 25)
53 print('Training accuracies per epoch:\n')
54 train_accuracy = trainLSTM(model, train_data_final_split, train_labels_final_split, 5, 0.00001)
55 print('LSTM Training Accuracy = ', train_accuracy)

```

☞ Training accuracies per epoch:

```

Epoch:  1  =  74.72 %
Epoch:  2  =  99.045 %
Epoch:  3  =  99.751 %
Epoch:  4  =  99.917 %
Epoch:  5  =  100.0 %

```

```

Training time =  228  seconds
LSTM Training Accuracy =  100.0

```

```

1  import time
2  def test_model(model, video_features, video_labels):
3      device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
4      start_time = time.time()
5      model.eval()
6      with torch.no_grad():
7          correct_count = 0

```

```

7         correct_count = 0
8         total_count = 0
9         for i in range(video_features.shape[0]):
10             video_feature = video_features[i]
11             label_feature = video_labels[i].type(torch.LongTensor)
12             label_feature = label_feature.view(1) - 1
13
14             preds = model(video_feature.cuda(device))
15             preds = preds.view(1, preds.shape[0])
16             _, pred_vals = torch.max(preds.data, 1)
17             total_count = total_count + label_feature.size(0)
18             correct_count = correct_count + (pred_vals == label_feature.cuda(device)).sum().item()
19
20         print('Test accuracy = ', round((correct_count / total_count) * 100, 3), '%')
21
22     print('Testing time: ', round(time.time() - start_time), ' seconds')
23     return (correct_count / total_count) * 100
24 test_accuracy = test_model(model, test_data_final_split, test_labels_final_split)
25 print('LSTM Testing Accuracy = ', test_accuracy)

```

```

➤ Test accuracy = 80.967 %
Testing time: 4 seconds
LSTM Testing Accuracy = 80.96740273396425

```

Reference:

<https://www.jessicayung.com/lstms-for-time-series-in-pytorch/>

```

1 train_for_linearSVC = train_data_final_split.view(train_data_final_split.shape[0], train_data_final_split.shape[1]* train_data_final_split.shape[2])
2 train_labels_for_linearSVC = train_labels_final_split.cpu().detach().numpy()
3 test_for_linearSVC = test_data_final_split.view(test_data_final_split.shape[0], test_data_final_split.shape[1]* test_data_final_split.shape[2])
4 test_labels_for_linearSVC = test_labels_final_split.cpu().detach().numpy()

```

```

1 from sklearn.svm import LinearSVC
2 from sklearn.metrics import accuracy_score
3 import time
4 start_time = time.time()
5 linearSVCClassifier = LinearSVC(C=0.0001, multi_class="ovr")
6 linearSVCClassifier.fit(train_for_linearSVC, train_labels_for_linearSVC)
7 pred = linearSVCClassifier.predict(test_for_linearSVC)
8 test_accuracy_SVC = accuracy_score(pred, test_labels_for_linearSVC)*100

```

```

8 test_accuracy_SVC = accuracy_score(pred, test_labels_for_linearSVC)*100
9 print('Test accuracy: ', test_accuracy_SVC, '%')
10 print('Testing time: ', round(time.time() - start_time), ' seconds')

```

```

↳ Test accuracy: 86.01472134595163 %
Testing time: 195 seconds

```

```

1 linearSVCClassifier = LinearSVC(C=0.0001, multi_class="ovr")
2 linearSVCClassifier.fit(train_for_linearSVC, train_labels_for_linearSVC)
3 pred = linearSVCClassifier.predict(train_for_linearSVC)
4 train_accuracy_SVC = accuracy_score(pred, train_labels_for_linearSVC)*100
5 print('Training accuracy: ', train_accuracy_SVC, '%')
6 print('Training time: ', round(time.time() - start_time), ' seconds')

```

```

↳ Training accuracy: 100.0 %
Training time: 388 seconds

```

▼ Problem 3. Evaluation

```

1 # \*write your codes for evaluation (You can use multiple cells, this is just a place holder)

```

▼ Print the train and test accuracy of your model - Printed Above

```

1 # Don't hardcode the train and test accuracy
2 print('Training accuracy for LSTM is %2.3f :' %(train_accuracy) )

```

```

↳ Training accuracy for LSTM is 100.000 :

```

```

1 print('Test accuracy for LSTM is %2.3f :' %(test_accuracy) )

```

```

↳ Test accuracy for LSTM is 80.967 :

```

▼ Print the train and test and test accuracy of SVM

```

1 # Don't hardcode the train and test accuracy
2 print('Training accuracy for SVC is %2.3f' %(train_accuracy SVC))

```

```
- print('Training accuracy for SVC is %.3f' %(test_accuracy_SVC))
```

```
➤ Training accuracy for SVC is 100.000
```

```
1 print('Test accuracy for SVC is %.3f' %(test_accuracy_SVC))
```

```
➤ Test accuracy for SVC is 86.015
```

Problem 4. Report

▼ Bonus

▼ Print the size of your training and test data

```
1 # Don't hardcode the shape of train and test data
2 print('Shape of training data is :', )
3 print('Shape of test/validation data is :', )
```

▼ Modelling and evaluation

```
1 #Write your code for modelling and evaluation
```

▼ Submission

Runnable source code in ipynb file and a pdf report are required.

The report should be of 3 to 4 pages describing what you have done and learned in this homework and report performance of your model. If you have tried multiple methods, please compare your results. If you are using any external code, please cite it in your report. Note that this homework is designed to help you explore and get familiar with the techniques. The final grading will be largely based on your prediction accuracy and the different methods you tried (different architectures and parameters).

Please indicate clearly in your report what model you have tried, what techniques you applied to improve the performance and report their accuracies. The report should be concise and include the highlights of your efforts. The naming convention for report is

Surname_Givenname_SBUID_report*.pdf

When submitting your .zip file through blackboard, please -- name your .zip file as **Surname_Givenname_SBUID_hw*.zip**.

This zip file should include:

```
Surname_Givenname_SBUID_hw*
| ---Surname_Givenname_SBUID_hw*.ipynb
| ---Surname_Givenname_SBUID_hw*.pdf
| ---Surname_Givenname_SBUID_report*.pdf
```

For instance, student Michael Jordan should submit a zip file named "Jordan_Michael_111134567_hw5.zip" for homework5 in this structure:

```
Jordan_Michael_111134567_hw5
| ---Jordan_Michael_111134567_hw5.ipynb
| ---Jordan_Michael_111134567_hw5.pdf
| ---Jordan_Michael_111134567_report*.pdf
```

The **Surname_Givenname_SBUID_hw*.pdf** should include a **google shared link**. To generate the **google shared link**, first create a folder named **Surname_Givenname_SBUID_hw*** in your Google Drive with your Stony Brook account.

Then right click this folder, click **Get shareable link**, in the People textfield, enter two TA's emails: bo.cao.1@stonybrook.edu and sayontan.ghosh@stonybrook.edu. Make sure that TAs who have the link **can edit, not just can view**, and also **uncheck** the **Notify people** box.

Colab has a good feature of version control, you should take advantage of this to save your work properly. However, the timestamp of the submission made in blackboard is the only one that we consider for grading. To be more specific, we will only grade the version of your code right before the timestamp of the submission made in blackboard.

You are encouraged to post and answer questions on Piazza. Based on the amount of email that we have received in past years, there might be delays in replying to personal emails. Please ask questions on Piazza and send emails only for personal issues.

Please ensure that your code will run on a JupyterLab environment, both locally and on a remote server. Please do your own work.

Drive Link: https://drive.google.com/drive/folders/1Wi5pmgQ_hBVxGJOBYf943Q5kc01ff6fl?usp=sharing

